

# Telecom Demo - Complete Startup & Execution Guide

---

## Pre-Flight Checklist

---

**CRITICAL:** Complete this checklist BEFORE attempting the demo to ensure 100% success rate.

### 1. Docker Desktop Status

```
# Check Docker is running
docker --version
docker ps

# Expected: Version output and running containers
# ❌ Error: "Cannot connect to Docker daemon" → Start Docker Desktop
```

### 2. Kubernetes Cluster Health

```
# Verify Kubernetes is enabled and running
kubectl get nodes

# Expected: STATUS = Ready
# ❌ Error: "connection refused" → Enable Kubernetes in Docker Desktop
```

### 3. Core Services Status

```
# Check all pods are running
kubectl get pods -n ecommerce-ai

# Expected: All pods show STATUS = Running
# ❌ Any pod showing "Pending/Error/CrashLoopBackOff" → Restart that service
```

---

## Complete Startup Sequence

---

Follow this sequence **exactly** - each step depends on the previous ones.

### STEP 1: Start Docker Desktop

```
# macOS: Open Docker Desktop application
open -a Docker
```

```
# Wait for Docker to fully start (green status icon)
# Verify with:
docker ps
```

**What this does:** Starts the container runtime that hosts Kubernetes and all services.

#### Common Issues:

- Docker Desktop not installed → Download from [docker.com](https://docker.com)
- "Docker Desktop starting..." → Wait 2-3 minutes for full startup
- Port conflicts → Close other applications using ports 7687, 9092

## STEP 2: Verify Kubernetes Services

```
# Check all services are running
kubectl get pods -n ecommerce-ai

# Should show:
# kafka-confluent-xxx          Running
# neo4j-release-0              Running
```

**What this does:** Confirms your database and messaging services are operational.

#### If services are not running:

```
# Restart Kafka
kubectl rollout restart deployment/kafka-confluent -n ecommerce-ai

# Restart Neo4j
kubectl rollout restart statefulset/neo4j-release -n ecommerce-ai

# Wait 2-3 minutes, then verify:
kubectl get pods -n ecommerce-ai
```

## STEP 3: Network Connectivity Health Check

```
# Test Neo4j connectivity
nc -zv localhost 7687
# Expected: "Connection to localhost port 7687 [tcp] succeeded!"

# Test Kafka connectivity
nc -zv localhost 9092
# Expected: "Connection to localhost port 9092 [tcp] succeeded!"

# Test Neo4j Browser (optional)
curl -s http://localhost:7474 | grep -q "neo4j" && echo "Neo4j Browser OK" || echo "Neo4j Browser failed"
```

**What this does:** Verifies network paths from your Mac to the containerized services.

### If connectivity fails:

```
# Check service endpoints
kubectl get svc -n ecommerce-ai

# Look for LoadBalancer services with localhost:
# neo4j-lb-neo4j          LoadBalancer    localhost    7474:xxx/TCP,7687:xxx/TCP
# kafka-confluent-service LoadBalancer    localhost    9092:xxx/TCP
```

## STEP 4: Database Validation 🗄️

```
# Navigate to project directory
cd /Users/timtadeo/Desktop/TelcoDemoDAIR

# Test Neo4j connection with credentials
python3 -c "
from neo4j import GraphDatabase
try:
    driver = GraphDatabase.driver('bolt://localhost:7687', auth=('neo4j', 'password'))
    with driver.session() as session:
        result = session.run('MATCH (t:Tower) RETURN count(t) as towers')
        count = result.single()['towers']
        print(f'✅ Neo4j OK: {count} towers found')
    driver.close()
except Exception as e:
    print(f'❌ Neo4j Failed: {e}')
"
```

**Expected Output:** ✅ Neo4j OK: 100 towers found

### If database is empty or fails:

```
# Re-seed the database
python3 seed_neo4j.py

# Expected output:
# ✅ Constraints and indexes created successfully
# ✅ Seeded 4 carriers
# ✅ Seeded 100 towers
```

## STEP 5: Test Individual Scripts 🛠️

```
# Test tower data generator (30-second test)
timeout 30s python3 TelcoTowerDataV14.py

# Expected output:
# ✅ Loaded 100 towers from Neo4j
# 🎨 Starting Tower Data Generator...
# ✅ Tower status sent: TWR_0001 to tower-status

# Test CDR generator (30-second test)
```

```
timeout 30s python3 TelcoCDRDataKafkaV11.py
```

```
# Expected output:
# ✅ Loaded 100 towers and 4 carriers from Neo4j
# 🚀 Starting Telecom CDR Generator...
# ✅ CDR sent: CDR_123456 (telecom-cdr)
```

**What this does:** Validates that your Python scripts can connect to both Neo4j and Kafka independently.

**If scripts fail:**

- **Neo4j errors:** Check password in scripts matches `password123`
- **Kafka errors:** Verify Kafka service is running and accessible
- **Import errors:** Run `pip install faker confluent-kafka neo4j flask`

## STEP 6: Start Flask Webhook Service 🌐

```
# Start the webhook listener
python3 run_demo.py

# Expected output:
# 🌐 Telecom Demo Webhook Listener Starting...
# 📡 Listening on port 5001
# * Running on http://127.0.0.1:5001
```

**What this does:** Creates the API endpoint that n8n will call to trigger your demo.

**Keep this terminal open** - you'll see real-time demo output here.

**If port 5001 is in use:**

```
# Find what's using the port
lsof -i :5001

# Kill the process or change port in run_demo.py
```

## STEP 7: Test Local Webhook 🔧

**Open a new terminal** and test:

```
# Test health endpoint
curl -X GET http://localhost:5001/webhook/health

# Expected response:
# {"status": "healthy", "timestamp": "2025-01-20T..."}

# Test demo trigger (1-minute test)
curl -X POST http://localhost:5001/webhook/start-demo \
  -H "Content-Type: application/json" \
```

```
-d '{"duration_minutes": 1}'

# Expected response:
# {"status": "success", "message": "Demo started successfully", ...}
```

**What this does:** Confirms your Flask service can receive requests and start the scripts.

**Watch the Flask terminal** - you should see both scripts start and generate data.

## STEP 8: Expose to Internet 🌐

**In another new terminal:**

```
# Start ngrok tunnel
ngrok http 5001

# Expected output:
# Forwarding https://xxxx-68-118-242-43.ngrok-free.app -> http://localhost:5001
```

**Copy the HTTPS URL** - you'll need this for n8n.

**Test public access:**

```
# Replace with your actual ngrok URL
curl -X GET https://your-ngrok-url.ngrok-free.app/webhook/health

# Expected: Same healthy response as local test
```

## STEP 9: Configure n8n Workflow ⚡

In your n8n Cloud interface:

1. **Update HTTP Request URL** with your ngrok URL:

```
https://your-ngrok-url.ngrok-free.app/webhook/start-demo
```

2. **Verify configuration:**

- Method: `POST`
- Headers: `Content-Type: application/json`
- Body Parameters: `duration_minutes: 2`

3. **Save workflow**

---

# Demo Execution & Monitoring

## Execute the Demo

1. Click **"Execute workflow"** in n8n
2. **Monitor Flask terminal** for real-time output:

```
📞 Demo start request received - Duration: 2 minutes
[TelcoTowerDataV14.py] ✅ Loaded 100 towers from Neo4j
[TelcoCDRDataKafkaV11.py] ✅ Loaded 100 towers and 4 carriers from Neo4j
[TelcoTowerDataV14.py] ✅ Tower status sent: TWR_0008 to tower-status
[TelcoCDRDataKafkaV11.py] ✅ CDR sent: CDR_782139 (telecom-cdr)
[TelcoCDRDataKafkaV11.py] 📁 Stored interaction #1 in Neo4j (Status: DROPPED)
```

## Success Indicators

- ✅ n8n shows green checkmark and success response
- ✅ Flask terminal shows both scripts running with data flow
- ✅ Scripts complete after specified duration
- ✅ No error messages in any terminal

## Expected Data Volume (2-minute demo)

- **Tower Messages:** ~60 (1 every 2 seconds)
- **CDR Messages:** ~120 (1 every 1 second)
- **Neo4j Records:** ~12 (every 10th CDR stored)

## Health Check Commands

Use these commands to diagnose issues during demo:

### Quick Service Check

```
# All-in-one health check
echo "=== Docker ===" && docker ps --format "table {{.Names}}\t{{.Status}}" | head -n 1
echo "=== Kubernetes ===" && kubectl get pods -n ecommerce-ai --no-headers | awk '{print $2}'
echo "=== Connectivity ===" && nc -zv localhost 7687 && nc -zv localhost 9092
echo "=== Flask ===" && curl -s http://localhost:5001/webhook/health | grep -o 'OK'
```

### Neo4j Data Check

```
# Verify recent data
python3 -c "
from neo4j import GraphDatabase
driver = GraphDatabase.driver('bolt://localhost:7687', auth=('neo4j', 'password123'))
with driver.session() as session:
    result = session.run('MATCH (t:Tower) RETURN t')
    print('Recent Towers:', [record['t'] for record in result])"
```

```

with driver.session() as session:
    towers = session.run('MATCH (t:Tower) RETURN count(t)').single()[0]
    recent = session.run('MATCH (i:Interaction) WHERE i.Timestamp >= datetime()
    print(f'Towers: {towers}, Recent Interactions: {recent}')
driver.close()

```

## Common Error Scenarios & Solutions

Error Symptom	Cause	Solution
Cannot connect to Docker daemon	Docker Desktop stopped	Start Docker Desktop application
connection refused to kubectl	Kubernetes disabled	Enable Kubernetes in Docker Desktop settings
ServiceUnavailable: localhost:7687	Neo4j service down	<code>kubectl rollout restart statefulset/neo4j-release -n ecommerce-ai</code>
ApiVersionRequest failed	Kafka service down	<code>kubectl rollout restart deployment/kafka-confluent -n ecommerce-ai</code>
The endpoint xxx.ngrok-free.app is offline	ngrok tunnel expired	Restart ngrok, update n8n URL
port 5001 already in use	Previous Flask instance running	<code>lsof -i :5001</code> then kill process
No towers found	Database not seeded	Run <code>python3 seed_neo4j.py</code>
Import errors	Missing Python packages	Run pip install command

## Pre-Demo Final Checklist

Complete this 2-minute checklist before EVERY demo execution:

```

# 1. Docker running?
docker ps > /dev/null && echo "✅ Docker OK" || echo "❌ Start Docker Desktop"

# 2. Services running?

```

```
kubectl get pods -n ecommerce-ai | grep -q "Running" && echo "✅ K8s OK" || echo "❌ K8s not running"

# 3. Connectivity working?
nc -zv localhost 7687 > /dev/null 2>&1 && nc -zv localhost 9092 > /dev/null 2>&1 && echo "✅ Connectivity OK" || echo "❌ Connectivity not working"

# 4. Flask responsive?
curl -s http://localhost:5001/webhook/health | grep -q "healthy" && echo "✅ Flask responsive" || echo "❌ Flask not responsive"

# 5. ngrok tunnel active?
curl -s https://your-ngrok-url.ngrok-free.app/webhook/health | grep -q "healthy" && echo "✅ ngrok tunnel active" || echo "❌ ngrok tunnel not active"
```

**Only proceed with n8n execution if ALL checks show ✅**

---

## Troubleshooting Quick Commands 🚨

---

**If demo fails mid-execution:**

```
# Stop all running processes
pkill -f "python3.*Telco"
pkill -f "run_demo.py"

# Restart everything fresh
kubectl rollout restart deployment/kafka-confluent -n ecommerce-ai
kubectl rollout restart statefulset/neo4j-release -n ecommerce-ai

# Wait 3 minutes, then restart from STEP 6
```

**Emergency reset:** If nothing works, restart Docker Desktop completely and begin from STEP 1.

---

**Following this guide ensures your demo will work perfectly on the first try every time!**

